

# Tutorial: Rejection Sampling

Matteo Fasiolo

# Sampling

Sampling is one of the most common tasks in statistical data analysis and in particular Monte Carlo simulations.

The task of sampling is to draw samples of a random variable (r.v.), given its probability density/mass function.

R provides lots of built-in samplers following the pattern:


```
rdistr(n, par1, par2, par3, ...)
```

which outputs a vector of  $n$  random variable from a distribution with parameters `par1`, `par2`, `par3`, ...

Example: sample from  $N(\mu, \sigma^2)$  with  $\mu = 0$  and  $\sigma = 2$ :


```
x <- rnorm(n = 1e4, mean = 0, sd = 2)
hist(x)
```

## Histogram of x



Sample from  $Unif(a, b)$  with  $a = -1$  and  $b = 1$ :

```
x <- runif(n = 1e4, min = -1, max = 1)
hist(x, prob = TRUE, main = "")
```




We use `prob = TRUE` to have probability densities on *y*-axis.

## Sampling for non-standard distributions

But, how do you sample from (e.g.) a triangular distribution?

Suppose we know its p.d.f.  $p(x) = 1 - |x - 1|, x \in [0, 2]$ ?



# Solution: Rejection Sampling

Basic idea:


1. sample from a distribution you know how to sample from (the **proposal distribution**);
2. “adjust” that sample so it becomes a sample from the distribution of interest (the **target distribution**).

Let  $p(x)$  be the p.d.f. of target and  $q(x)$  be the p.d.f. of proposal.

Let  $M^* > 1$  be smallest constant such that:

$$p(x) \leq M^* q(x) \text{ for any } x \text{ s.t. } p(x) > 0.$$

Sounds confusing? We just need to  $Mq(x)$  to “cover”  $p(x)$ ...



NB: you can choose  $M > M^*$ , but  $p(x) \leq Mq(x)$  is necessary.

# Algorithmic Implementation

Objective: get a sample of  $n_{\max}$  independent random variables (r.v.s) from target.

Algorithm:

1. Draw a r.v.  $x$  from  $q(x)$  and a uniform sample  $u \sim U(0, 1)$
2. If  $u < \frac{p(x)}{M^*q(x)}$ , accept sample  $x$  and store it.
3. If number of accepted samples so far equals to  $n_{\max}$ , quit.  
Otherwise, go back to step 1.

Here we will use  $Unif(a = 0, b = 2)$  as proposal distribution  $q(x)$ .



When accepting a sample in step 2, we need to add it to the vector of accepted samples so far.

Two ways of doing it:

```
a <- c(1,2,3,4)
```

```
a[5] <- 5
```

```
print(a)
```

```
[1] 1 2 3 4 5
```

```
# OR
```

```
a <- c(a, 6)
```

```
print(a)
```

```
[1] 1 2 3 4 5 6
```

# Code Skeleton

```
# how many samples do we want?
n_max = 1000

p_target <- function(x){
  # TODO: the PDF of distribution, from which
  # you want to sample.
}

acc <- c() # create an empty vector
n <- 0 # how many samples have we already obtained?
M <- NA # TODO you should be able to work it out!
while(n < n_max){
  # TODO
}
hist(acc) #plot the histogram of accepted samples
```

## Questions:

1. Complete the above code skeleton according to the rejection sampling algorithm.
2. What output is expected when the algorithm is implemented correctly?
3. What happens if you increase  $M$  above  $M^*$ ? Is the algorithm still correct (use visualisation to assess this)? Does the code become faster or slower?
4. What happens if you decrease  $M$  below minimum required value?

Challenge: try using  $N(\mu = 1, \sigma^2 = 0.5^2)$  as proposal:

- ▶ Computing optimal  $M^*$  might be hard, but you just need  $M$  big enough.
- ▶ To compute it you need to know the Gaussian p.d.f. (see `?dnorm`).
- ▶ Is this more efficient than using a uniform proposal?