

# Data frames: a data structure for statistical analysis

Matteo Fasiolo

## Creating a Scoring Table

Imagine we wanted to create a “table” which stores data of different types:

student	number,	name,	staff,	score
1		teo,	T,	70
2		jack,	T,	80
3		emma,	F,	70
...				

We cannot use `matrix` as it only allows single type data.

Lists can be used to store data of different types, can we use that?

```
stu_no <- c(1L, 2L, 3L)
name <- c("teo", "jack", "emma")
staff <- c(T, T, F)
score <- c(70, 80, 70)
score_table <- list(no = stu_no, name = name,
                    staff = staff, score = score)
```

This works:

```
score_table$name
```

```
[1] "teo" "jack" "emma"
```

```
score_table[[2]]
```

```
[1] "teo" "jack" "emma"
```

```
score_table$name[1]
```

```
[1] "teo"
```

But recall the data

student	number	name	staff	score
1		teo,	T,	70
2		jack,	T,	80
3		emma,	F,	70
...				

To extract all information on Emma we must do:

```
score_table[[1]][3]
```

```
[1] 3
```

```
score_table[[2]][3]
```

```
[1] "emma"
```

```
score_table[[3]][3]
```

```
[1] FALSE
```

(Exercise: can you do this more compactly with `lapply?`)

# Data Frames

Point is that list does not directly represent tabular format.

A data.frame does:

```
score_table <- data.frame(stu_no, name, staff, score)
score_table
```

	stu_no	name	staff	score
1	1	teo	TRUE	70
2	2	jack	TRUE	80
3	3	emma	FALSE	70

Extract info about Emma:

```
score_table[3, ]
```

	stu_no	name	staff	score
3	3	emma	FALSE	70

So data.frame can be accesses like a matrix:

```
score_table[3, 1]
```

```
[1] 3
```

```
score_table[1:2, ]
```

	stu_no	name	staff	score
1	1	teo	TRUE	70
2	2	jack	TRUE	80

We can also perform numeric operations on numeric columns:

```
score_table[, c(1, 4)] * 2 # Try score_table * 2
```

	stu_no	score
1	2	140
2	4	160
3	6	140

Can also access a data.frame like a list:

```
score_table$name
```

```
[1] "teo" "jack" "emma"
```

```
score_table[[2]]
```

```
[1] "teo" "jack" "emma"
```

```
score_table[c(1, 2)]
```

```
  stu_no name
1      1  teo
2      2 jack
3      3 emma
```

If can also mix matrix-like and list-like indexing:

```
score_table[score_table$name == "emma", ]
```

```
  stu_no name staff score
3      3 emma FALSE    70
```

Using list-like access we add/remove columns (variables):

```
score_table$age <- c(23, 23, 21)
score_table
```

	stu_no	name	staff	score	age
1	1	teo	TRUE	70	23
2	2	jack	TRUE	80	23
3	3	emma	FALSE	70	21

```
score_table$age <- NULL
```

Using matrix-like access we can remove or add:

```
( score_table <- score_table[-1, ] )
```

	stu_no	name	staff	score
2	2	jack	TRUE	80
3	3	emma	FALSE	70

```
score_table[3, ] <- list(1, "Jacob", TRUE, 60) #Try c()
```



# Summarizing Data Frame

You can also get a summary of your data frame:

```
summary(score_table)
```

	stu_no	name	staff
Min.	:1.0	Length:3	Mode :logical
1st Qu.:	1.5	Class :character	FALSE:1
Median	:2.0	Mode :character	TRUE :2
Mean	:2.0		
3rd Qu.:	2.5		
Max.	:3.0		

data.frames are a basic input for many modelling functions: `lm`, `glm`, `gam`, ...

## Case Study: Predicting MNIST Digits.

Let us revisit the TB1 project and how data frame can be useful in this specific application.

We have the following matrices

- ▶ Images:  $X \in \mathbb{R}^{60000 \times 784}$  ( $\sqrt{784} = 28$ )
- ▶ Labels:  $Y \in \{1, 2, 3, 4, 5, 6, 7, 8, 9, 0\}^{60000 \times 1}$
- ▶ Testing Images:  $TX \in \mathbb{R}^{10000 \times 784}$
- ▶ Testing Labels:  $TY \in \{1, 2, 3, 4, 5, 6, 7, 8, 9, 0\}^{10000 \times 1}$

We can directly use these matrices within our program (as you did in the TB1 project).

However, by using data frame, we can organize them in a more meaningful way, making our program easily understandable.

Let us download the data

```
library(dslabs)
# Need internet access!
# Do save(file = "mnist.RData", mnist)
# Then load("mnist.RData")
mnist <- read_mnist()
str(mnist)
```

List of 2

\$ train:List of 2

..\$ images: int [1:60000, 1:784] 0 0 0 0 0 0 0 0 0 0 0 ...

..\$ labels: int [1:60000] 5 0 4 1 9 2 1 3 1 4 ...

\$ test :List of 2

..\$ images: int [1:10000, 1:784] 0 0 0 0 0 0 0 0 0 0 0 ...

..\$ labels: int [1:10000] 7 2 1 0 4 1 4 9 5 9 ...

We rearrange the data and subsample:

```
sub <- 1:1000
train <- list("images" = mnist$train$images[sub, ],
             "labels" = mnist$train$labels[sub],
             "istest" = rep(FALSE, length(sub)))
test <- list("images" = mnist$test$images[sub, ],
            "labels" = mnist$test$labels[sub],
            "istest" = rep(TRUE, length(sub)))
```

```
dat_mnist <- data.frame(  
  "images" = I( rbind(train$images, test$images) ),  
  "labels" = c(train$labels, test$labels),  
  "istest" = c(train$istest, test$istest))
```

```
dim(dat_mnist)
```

```
[1] 2000    3
```

```
dim(dat_mnist$images)
```

```
[1] 2000  784
```

```
str(dat_mnist)
```

```
'data.frame':   2000 obs. of  3 variables:  
 $ images: 'AsIs' int [1:2000, 1:784] 0 0 0 0 0 0 0 0 0 0 0 ...  
 $ labels: int 5 0 4 1 9 2 1 3 1 4 ...  
 $ istest: logi FALSE FALSE FALSE FALSE FALSE FALSE ...
```

The I function indicates we want the list elements stored in the data frame **as is**.

Without using I function, R will break the images matrix into its columns.

```
dat_wrong <- data.frame(  
  "images" = rbind(train$images, test$images),  
  "labels" = c(train$labels, test$labels),  
  "istest" = c(train$istest, test$istest))
```

```
dim(dat_wrong)
```

```
[1] 2000  786
```

```
str(dat_wrong)  
'data.frame':    2000 obs. of  786 variables:  
 $ images.1  : int  0 0 0 0 0 0 0 0 0 0 0 ...  
 $ images.2  : int  0 0 0 0 0 0 0 0 0 0 0 ...  
 $ images.3  : int  0 0 0 0 0 0 0 0 0 0 0 ...
```

## Visualizing Images

```
library(imager)

img <- matrix(dat_mnist$images[1, ], nrow = 28)
plot(as.cimg(img), axes = FALSE)
title(dat_mnist$labels[1])
```

5



For our convenience, let us add the image objects for each digit as a new column to the data frame, so we can easily plot them:

```
cvrt <- function(imgdata){  
  img <- imgdata  
  img <- matrix(img, nrow = 28, ncol = 28)  
  return(as.cimg(img))  
}  
  
dat_mnist$pics <- apply(dat_mnist$image, 1, cvrt,  
                        simplify = FALSE)  
dim(dat_mnist) # now we have 4 columns.
```

```
[1] 2000    4
```

```
str(dat_mnist$pics[1:3])
```

List of 3

```
$ : 'cimg' num [1:28, 1:28, 1, 1] 0 0 0 0 0 0 0 0 0 0 0 0 ...  
$ : 'cimg' num [1:28, 1:28, 1, 1] 0 0 0 0 0 0 0 0 0 0 0 0 ...  
$ : 'cimg' num [1:28, 1:28, 1, 1] 0 0 0 0 0 0 0 0 0 0 0 0 ...
```



```
str(dat_mnist)
```

The columns are a matrix, two vectors and a list of graphical arrays!

Now each row identifies a digit.

To plot one we do:

```
plot(dat_mnist$pics[[5]], axes = FALSE,  
     main = dat_mnist$labels[5])
```

**9**



## Find Neighbours

First, let us write a function produces the 20 nearest neighbours who have the smallest distances to a test image

```
# returns indices of 20 nearest neighbours
find_nei <- function(a, dat){
  train <- dat[dat$istest == FALSE, ]
  n <- nrow(train$images)
  dist <- numeric(n)
  for(i in 1:n){
    dist[i] <- sqrt(sum((a-train$images[i, ])^2))
  }
  return(order(dist)[1:20])
}
```

Where note that

```
order(c(3.5, 1.1, 2))
```

```
[1] 2 3 1
```

Test our function

```
kk <- 1000 + 2 # 2nd digit in test set  
nei <- find_nei(dat_mnist$images[kk, ], dat = dat_mnist)  
dat_mnist$labels[nei]
```

```
[1] 2 2 8 3 2 6 3 3 2 2 2 2 6 1 6 3 1 6 2 1
```

```
plot(dat_mnist$pics[[kk]], axes = FALSE,  
      main = dat_mnist$labels[kk])
```

**2**



# Making Predictions

We want to predict the label based on the majority of the neighbours' labels. For example, here:

```
dat_mnist$labels[nei]
```

```
[1] 2 2 8 3 2 6 3 3 2 2 2 2 6 1 6 3 1 6 2 1
```

our prediction should be 2.

A useful command here is

```
( tab <- table(dat_mnist$labels[nei]) )
```

```
1 2 3 6 8
3 8 4 4 1
```

```
names( tab )
```

```
[1] "1" "2" "3" "6" "8"
```

We want the label corresponding to the largest count:

```
names(tab)[tab == max(tab)]
```

```
[1] "2"
```

So our prediction function is:

```
predict_lab <- function(a, dat){  
  
  nei <- find_nei(a, dat)  
  
  nei_labs <- dat$labels[nei]  
  
  tab <- table(nei_labs)  
  
  pred <- names(tab)[tab == max(tab)][1]  
  
  return(pred)  
  
}
```

Check if it works:

```
predict_lab(dat_mnist$images[kk, ], dat = dat_mnist)
```

```
[1] "2"
```

Apply predict function to all images.

```
dat_mnist$preds <- apply(dat_mnist$images, 1,  
                        predict_lab, dat = dat_mnist)
```

Exercise: this takes a while, what is the order of complexity of the computation above?

Hint: how many pairs of images are we comparing and what is the complexity of each comparison?

Add error columns:

```
dat_mnist$error <- as.integer(dat_mnist$preds) -  
                    as.integer(dat_mnist$labels)
```

```
train_err <- dat_mnist$error[dat_mnist$is_test == FALSE]  
test_err <- dat_mnist$error[dat_mnist$is_test == TRUE]  
table(train_err)
```

train\_err

-9	-8	-7	-6	-5	-4	-3	-2	-1	0	1	2	3	4	5
2	2	12	13	13	2	13	24	20	849	7	8	3	5	19

```
table(test_err)
```

test\_err

-8	-7	-6	-5	-4	-3	-2	-1	0	1	2	3	4	5	6
7	8	14	15	9	9	39	48	765	12	14	1	10	41	7



Mean absolute errors

```
mean( abs( dat_mnist$error[dat_mnist$isitest == FALSE] ) )
```

```
[1] 0.575
```

```
mean( abs( dat_mnist$error[dat_mnist$isitest == TRUE] ) )
```

```
[1] 0.799
```

We do better on the train than on the testing set, why?

Is there a (simple) solution to (nearly) equalise the performance?

# Conclusion

A `data.frame` combines data of different types together, forms “datasets”.

- ▶ Datasets are organized as “observations” and “columns (variables)”.
- ▶ a `data.frame` can be accessed like a matrix or like a list.

A `data.frame` provides a clean and unified way of managing and accessing datasets in our program.

- ▶ Instead of managing several vectors, we only need to manage one data frame for each dataset.